



```
In [32]: import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
(x_train, y_train), (x_test, y_test) = mnist.load_data()
print(x_train.shape, y_train.shape)

(60000, 28, 28) (60000,)
```

```
In [33]: x_train= x_train.reshape(x_train.shape[0],28,28,1)
x_test= x_test.reshape(x_test.shape[0],28,28,1)
input_shape=(28,28,1)
y_train=keras.utils.to_categorical(y_train)#, num_classes=)
y_test=keras.utils.to_categorical(y_test)#, num_classes=)
x_train= x_train.astype('float32')
x_test= x_test.astype('float32')
x_train /= 255
x_test /= 255
```

```
In [45]: batch_size=64
num_classes=10
epochs=10
def build_model(optimizer):
```



```
model = Sequential()
model.add(Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy, optimizer= optimizer, metrics=['accuracy'])
return model
```

```
In [47]: optimizers = ['Adadelta', 'Adagrad', 'Adam', 'RMSprop', 'SGD']

for i in optimizers:

    model = build_model(i)

    hist=model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(
ss: 0.1615 - val_accuracy: 0.9538
Epoch 5/10
938/938 [=====] - 29s 31ms/step - loss: 0.2382 - accuracy: 0.9289 - val_lo
ss: 0.1437 - val_accuracy: 0.9560
Epoch 6/10
938/938 [=====] - 29s 30ms/step - loss: 0.2204 - accuracy: 0.9344 - val_lo
ss: 0.1314 - val_accuracy: 0.9620
Epoch 7/10
```

jupyter Assignment 5 DL Last Checkpoint: a minute ago (autosaved) Python 3 (ipykernel) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel)


```
938/938 [=====] - 29s 31ms/step - loss: 0.2067 - accuracy: 0.9378 - val_lo
ss: 0.1214 - val_accuracy: 0.9643
Epoch 8/10
938/938 [=====] - 30s 32ms/step - loss: 0.1956 - accuracy: 0.9412 - val_lo
ss: 0.1146 - val_accuracy: 0.9652
Epoch 9/10
938/938 [=====] - 29s 31ms/step - loss: 0.1833 - accuracy: 0.9453 - val_lo
ss: 0.1084 - val_accuracy: 0.9660
Epoch 10/10
938/938 [=====] - 28s 30ms/step - loss: 0.1746 - accuracy: 0.9467 - val_lo
ss: 0.1032 - val_accuracy: 0.9680
```

```
In [50]: import matplotlib.pyplot as plt

# List of optimizers to compare
optimizers = ['Adadelata', 'Adagrad', 'Adam', 'RMSprop', 'SGD']

# Loss for each optimizer (replace with your actual Loss values)
loss_values = [1.1345, 0.3577, 0.0267, 0.0360, 0.1746] # Replace with actual Loss values

# Plotting the graph
plt.figure(figsize=(10, 6))
plt.bar(optimizers, loss_values, color=['blue', 'green', 'red', 'purple', 'orange'])
plt.xlabel('Optimizers')
plt.ylabel('Training Loss')
plt.title('Training Loss for Different Optimizers')
plt.ylim(0, max(loss_values) + 0.05) # Adjust the y-axis limits if needed
plt.show()
```



## Conclusion:

Through this experiment we have successfully developed a model to test the accuracy, loss and convergence time of various optimizers such as stochastic gradient descent, gradient descent, adam, adagrad, rmsprop and nadam.